

CS289A_HW06_NeuralNet

April 14, 2017

1 Neural Network Constructor

This notebook works through the construction of a two-layer neural network. (It will eventually be imported into its own python module)

```
In [108]: %load_ext autoreload
```

The autoreload extension is already loaded. To reload it, use:
%reload_ext autoreload

```
In [109]: %autoreload 2
```

```
In [110]: import numpy as np
import HW06_utils as ut

import gradients as grad
import activationfns as af
```

```
In [462]: BASE_DIR = '/Users/mitch/Documents/Cal/2_2017_Spring/COMPSCI 289A - Intro
traindata = np.loadtxt(BASE_DIR+'Data/letters_traindata.csv', dtype=float)
trainlabels = np.array([np.loadtxt(BASE_DIR+'Data/letters_trainlabels.csv', dtype=int)
```

```
In [365]: class NeuralNet:
    """
    Train and store a neural network, based on supplied training data.
    Use this network to predict classifications.
    """

    def __init__(self, nlayers=3, unitsperlayer=None, actfns=[af.sigmoid, af.tanh]):
    """
    Initialize the neural network
    - nlayers: the number of layers in the neural network (including input and output layers)
    - unitsperlayer: a list specifying (in order) the number of units in each layer
    - actfns: a list specifying (in order) the activation functions for each layer
    - Gradient: a class providing optimized gradient calculation for the specified activation functions
    - verbose: a boolean for descriptive output
```

```

"""
if unitsperlayer == None:
    unitsperlayer = 3*np.ones(nlayers)
elif nlayers == len(unitsperlayer)+1:
    self.nlayers = nlayers-1
    self.unitsperlayer = unitsperlayer
elif nlayers > len(unitsperlayer)+1:
    print('ERROR: The number of units per layer were not given for')
elif nlayers < len(unitsperlayer)+1:
    print('ERROR: More layers were given units than were specified')
if nlayers == len(actfns)+1:
    self.actfns = actfns
elif nlayers > len(actfns)+1:
    print('ERROR: The activation function was not given for at least')
elif nlayers < len(actfns)+1:
    print('ERROR: More activation functions were provided than specified')
if Gradients == None:
    print('ERROR: A gradient generator class must be included.')
self.gradients = Gradients
self.weight_matrices = []

def initialize_weights(self, shape, mu=0, var=1):
    """
    Initialize weight matrix from normal distribution.
    - shape: tuple specifying desired shape of weight matrix
    - mu:    mean value of normal distribution
    - var:   variance of normal distribution
    """
    weight_matrix = np.random.normal(loc=mu, scale=np.sqrt(var), size=shape)

    return weight_matrix

def weight_matrix_shape(self, layer_n, nfeatures):
    """
    Create weight matrix with the proper number of rows and columns
    - n:          the layer which will employ an activation function on the
                  product of the weight matrix and values
    - nfeatures: an integer specifying the number of features in the
    """
    if layer_n != 0 and layer_n != range(self.nlayers)[-1]:
        WM_nrows = self.unitsperlayer[layer_n]-1
        WM_ncols = self.unitsperlayer[layer_n-1]
    elif layer_n == 0:
        WM_nrows = self.unitsperlayer[layer_n]-1
        WM_ncols = nfeatures
    else:

```

```

        WM_nrows = self.unitsperlayer[layer_n]
        WM_ncols = self.unitsperlayer[layer_n-1]
    return WM_nrows,WM_ncols

def forward(self,data):
    """
    Perform forward pass through neural network by multiplying data by weight matrices
    and enforcing a nonlinear activation function for each layer.
    - data: Nxd numpy array with N sample points and d features
    - weightmatrices: ordered list of sequential weight matrices corresponding to each layer
    - actfns: ordered list of sequential activation functions (functions are defined in activationfuncs.py)
    Returns layeroutputs, a list of the outputs from each layer. The final output is a CxN numpy array with hypotheses for each sample N_i being 1 or 0
    """
    H = data.T
    layeroutputs = []
    for i in range(self.nlayers):
        W = self.weight_matrices[i]
        actfn = self.actfns[i]
        H = actfn(np.dot(W,H))
        # If the layer is not the output layer, add a fictitious unit of 1
        if i != self.nlayers-1:
            fictu = np.array([np.ones_like(H[0])])
            H = np.concatenate((H,fictu),axis=0)
        layeroutputs.append(H)
    return layeroutputs

def backward(self,layeroutputs,labelrange,gradients=None):
    """
    Perform backward pass through neural network by computing gradients of the loss function with respect to the input weight matrices with respect to the loss function comparing predicted values to true values. Classes for gradients are provided in gradients.py (a unique gradient class is required for neural networks with different numbers of layers and/or different activation functions)
    """
    if gradients == None:
        Gradients = self.gradients
        gradients = Gradients.calculate(self.weight_matrices,layeroutputs)

    return gradients

def classify_outputs(self,finaloutputs):
    """
    Convert final outputs into classifications

```

```

- finaloutputs: a CxN numpy array with hypotheses for each sample
                  class C_j.
Returns a 1D, length-N array with values corresponding to point c
"""
if len(finaloutputs) == 1:
    classifications = np.around(finaloutputs[0]).astype(int)
if len(finaloutputs) > 1:
    # Add one for 1-indexing in classification labels
    classifications = (np.argmax(finaloutputs, axis=0) + np.ones(len(
return classifications

def stoch_grad_descent_prep(self, layeroutputs, wrongclass):
    """
    For stochastic gradient descent, choose one misclassified point i
    (index i) for performing backprop algorithm and reduce datasets a
    - layeroutputs: a list of the outputs from each layer
    - predictions: 1D, length-N numpy array with predictions for the
    - labels:      1D, length-N numpy array with true labels for the
    """
    diffclass = True
    tested_i = []

    while counter < len(wrongclass):
        i = np.random.randint(len(labels))
        if i not in tested_i:
            tested_i.append(i)
        if len(tested_i) == len(labels):
            print('All points classified correctly')
            return True, True
        if labels[i] != classifications[i]:
            # Improperly classified point, so use it for gradient des
            return layeroutputs_i, i

def train(self, data, labels, epsilon=0.1):
    """
    Train the neural network on input data
    - data:    Nxd numpy array with N sample points and d features
    - labels: 1D, length-N numpy array with labels for the N sample p
    """
    # Ensure labels are integers and that data and labels are the same
    labels = labels.astype(int)
    if len(data) != len(labels):
        print('ERROR: Data and labels must be the same length.')

    # Add fictitious unit for bias terms
    fictu = np.array([np.ones(len(data))]).T

```

```

data = np.concatenate((data,fictu),axis=1)

# Initialize Weights
nfeatures = len(data[0])
for layer_n in range(self.nlayers):
    WM_nrows,WM_ncols = self.weight_matrix_shape(layer_n,nfeatures)
    # Variance of weight matrix determined by fan-in (eta), the n
    # (or the number of data features when initializing the first
    eta = WM_ncols
    weight_matrix = self.initialize_weights((WM_nrows,WM_ncols),n
    self.weight_matrices.append(weight_matrix)

# Begin loop
layeroutputs = self.forward(data)
classifications = self.classify_outputs(layeroutputs[-1])
trainAccs = [ut.score_accuracy(classifications,labels)]
epochcounter = 0
while epochcounter < 20:
    # Stochastic gradient descent: Loop over points randomly, one
    # (Execute gradient class overhead before beginning)
    self.gradients.prepare(data,labels,self.unitsperlayer[-1])

    for datapoint_i in range(len(data)):
        X_i = np.array([data[datapoint_i]])
        layeroutput_i = self.forward(X_i)

        gradients = self.backward(layeroutput_i,[datapoint_i,data

        for n in range(self.nlayers):
            self.weight_matrices[n]=self.weight_matrices[n]-epsil

    layeroutputs = self.forward(data)
    classifications = self.classify_outputs(layeroutputs[-1])
    trainAcc = ut.score_accuracy(classifications,labels)
    trainAccs.append(trainAcc)

    epochcounter+=1
    print('%.2f%% accuracy after %i epoch(s).' %(100*trainAcc,epochcounter))
    DL = np.concatenate((data,labels),axis=1)
    np.random.shuffle(DL)
    data = DL[:, :-1]
    labels = np.array([DL[:, -1]]).T
    epsilon *=0.75
print(trainAccs)

def predict(self,testdata):
    """

```

```

Predict classifications for unlabeled data points using the previously
trained neural network.
- testdata: Nxd numpy array with N sample points and d features
  *Note, dimension d must match that used for the data
Returns a 1D, length-N numpy array of predictions (one prediction per
point)
"""

# Add fictitious unit to input to match dimensions
fictu = np.array([np.ones(len(testdata))]).T
testdata = np.concatenate((testdata, fictu), axis=1)

npoints = len(testdata)
predictions = np.empty(npoints)
layeroutputs = self.forward(testdata)
predictions = self.classify_outputs(layeroutputs[-1])

return predictions.astype(int)

```

1.0.1 Implementation

```
In [366]: classifier = NeuralNet(nlayers=3, unitsperlayer=[201, 26], actfns=[af.tanh, af.relu])
```

```
In [367]: classifier.train(traindata, trainlabels)
```

```

73.91% accuracy after 1 epoch(s).
79.92% accuracy after 2 epoch(s).
83.10% accuracy after 3 epoch(s).
83.48% accuracy after 4 epoch(s).
85.09% accuracy after 5 epoch(s).
86.20% accuracy after 6 epoch(s).
86.69% accuracy after 7 epoch(s).
87.58% accuracy after 8 epoch(s).
88.06% accuracy after 9 epoch(s).
88.04% accuracy after 10 epoch(s).
88.44% accuracy after 11 epoch(s).
88.55% accuracy after 12 epoch(s).
88.64% accuracy after 13 epoch(s).
88.75% accuracy after 14 epoch(s).
88.79% accuracy after 15 epoch(s).
88.82% accuracy after 16 epoch(s).
88.91% accuracy after 17 epoch(s).
88.91% accuracy after 18 epoch(s).
88.93% accuracy after 19 epoch(s).
88.94% accuracy after 20 epoch(s).
[0.03782051282051282, 0.7390725160256411, 0.7991686698717949, 0.8310196314102564, 0.8310196314102564, 0.8310196314102564, 0.8310196314102564, 0.8310196314102564, 0.8310196314102564, 0.8310196314102564, 0.8310196314102564, 0.8310196314102564, 0.8310196314102564, 0.8310196314102564, 0.8310196314102564, 0.8310196314102564, 0.8310196314102564, 0.8310196314102564, 0.8310196314102564, 0.8310196314102564]

```

```
In [445]: valdata = np.loadtxt(BASE_DIR+'Data/letters_valdata.csv', dtype=float, de
        vallabels = np.array([np.loadtxt(BASE_DIR+'Data/letters_vallabels.csv', c


In [446]: valpredictions = classifier.predict(valdata)

In [447]: valAcc = ut.score_accuracy(valpredictions,vallabels)
        print(valAcc)


0.8662660256410256
```

```
In [448]: import matplotlib.pyplot as plt
```

```
In [472]: for num in range(40,60):
        plt.imshow(valdata[num,:784].reshape((28,28)))
        plt.show()
        alphabet = 'abcdefghijklmnopqrstuvwxyz'
        print('Guess: \t',alphabet[valpredictions[num]-1])
        print('Actually:\t',alphabet[vallabels[num,0]-1])
```




```
Guess:          e
Actually:       e
```




Guess:
Actually:


g
g




Guess: j
Actually: j




Guess: u
Actually: u



Guess: u
Actually: m




Guess: q
Actually: q




Guess:
Actually:

b
b




Guess:
Actually:

w
w




Guess:
Actually:


v
v



Guess: v
Actually: v




Guess: s
Actually: s




Guess:
Actually:


1
y




Guess: j
Actually: j




Guess: d
Actually: a



Guess: q
Actually: q




Guess: x
Actually: x




Guess:
Actually:

o
q




Guess:
Actually:

d
d



Guess:
Actually:

1
1




```
Guess:          x
Actually:       x
```

```
In [465]: testdata = np.loadtxt(BASE_DIR+'Data/letters_testdata.csv', dtype=float,
```

```
In [466]: testpredictions = classifier.predict(testdata)
```


```
In [467]: plt.imshow(testdata[0,:784].reshape((28,28)))
           plt.show()
```




```
In [468]: def make_kaggle(predictions,kagglecsvfilename,indexing=0):
           # Use this optimal classifier on the test data
           if indexing == 0:
               ids = np.arange(len(predictions))
           elif indexing == 1:
               ids = np.arange(1,len(predictions)+1)
           predictions_csv = np.concatenate(([ids],[predictions]),axis=0).T
           np.savetxt(kagglecsvfilename,predictions_csv,fmt='%i',delimiter=',',header='')
           return predictions_csv
```

```
In [469]: kagglepreds = make_kaggle(testpredictions,BASE_DIR+'Kaggle/neuralnet_mneg
```


```
In [461]: for num in range(20):  
          plt.imshow(testdata[num,:784].reshape((28,28)))  
          plt.show()  
          alphabet = 'abcdefghijklmnopqrstuvwxyz'  
          print('Guess:\t',alphabet[testpredictions[num]-1])
```




Guess: q




Guess: 1




Guess: i




Guess: x




Guess: p




Guess: v



Guess: e




Guess: n




Guess:

o




Guess:


b




Guess: 1




Guess: a




Guess: z




Guess: x




Guess: q




Guess: q




Guess: q



Guess: i



Guess: n



Guess: k

In []: